

Performance testing focused on API latency and server load. We encountered "429 Quota Exceeded" and "504 Deadline Exceeded" errors due to API rate limits. This was successfully mitigated by implementing a strict 10-second timeout mechanism in the Python code and optimizing the application to use the high-speed `gemini-3.5-flash-lite` model.